

Design and Analysis of Algorithm

Divide-and-Conquer (II)

- 1 Fast Power/Exponentiation
- 2 Integer Multiplication
- 3 Matrix Multiplication
- 4 Polynomial Multiplication

Outline

1 Fast Power/Exponentiation

2 Integer Multiplication

3 Matrix Multiplication

4 Polynomial Multiplication

Fast Power/Exponentiation Problem

Input. $a \in \mathbb{R}$, $n \in \mathbb{N}$

Output. a^n

Naive algorithm. Sequential multiplication

$$a^n = \underbrace{a \cdot a \cdot \dots \cdot a}_{n} \cdot a$$

#(multiplication) = $n - 1$

Divide-and-Conquer: Divide

n is even

$$\underbrace{a \dots a}_{n/2} \mid \underbrace{a \dots a}_{n/2}$$

n is odd

$$\underbrace{a \dots a}_{(n-1)/2} \mid \underbrace{a \dots a}_{(n-1)/2} \mid a$$

$$a^n = \begin{cases} a^{n/2} \times a^{n/2} & n \text{ is even} \\ a^{(n-1)/2} \times a^{(n-1)/2} \times a & n \text{ is odd} \end{cases}$$

Basic operation. multiplication

- size of subproblem: smaller than $n/2$
- two subproblems (with size roughly $n/2$) are identical, only need computing once

$$W(n) = W(n/2) + \Theta(1)$$

master theorem (case 1) $\Rightarrow W(n) = \Theta(\log n)$

How to realize this algorithm? recursion vs. iteration

A Recursive Approach

Algorithm 1: Power(a, n): $a^n = (a^{-1})^{-n}$

- 1: **if** $n < 0$ **then return** Power($1/a, -n$); //handle negative integer exponent
 - 2: **if** $n = 0$ **then return** 1;
 - 3: **if** $n = 1$ **then return** a ;
 - 4: **if** n is even **then return** Power($a^2, n/2$);
 - 5: **if** n is odd **then return** $a \times$ Power($a^2, (n - 1)/2$);
-

Naive implementation: $x \leftarrow \text{Power}(a, n/2)$, return $x \times x$.

- has the same effect as above: cause we have to compute a^2

An Iterative Approach

$$y = a^n = a^{\sum_{i=0}^k b_k 2^k} = \prod_{i=0}^k (a^{2^k})^{b_k}$$

Algorithm 2: Square-and-Multiply(a, n)

```
1:  $(b_k, b_{k-1}, \dots, b_1, b_0) \leftarrow \text{BinaryDecomposition}(n);$ 
2:  $y \leftarrow 1;$ 
3:  $power \leftarrow a;$ 
4: for  $i = 0$  to  $k$  do
5:   if  $b_i = 1$  then  $y \leftarrow y \times power$ ; //add
6:    $power \leftarrow power \times power$                                 // double
7: end
8: return  $y$ 
```

Also known as binary exponentiation

Naturally extend to additive semigroups: **double-and-add**

Application of Power Algorithm

Fibonacci sequence: 1, 1, 2, 3, 5, 8, 13, 21, ...

add $F_0 = 0$, we obtain:

0, 1, 1, 2, 3, 5, 8, 13, 21, ...

Problem. Given initial values $F_0 = 0$, $F_1 = 1$, compute F_n

Naive algorithm. From F_0, F_1, \dots , repeatedly compute

$$F_n = F_{n-1} + F_{n-2}$$

Complexity. sequential addition: $\Theta(n)$

Properties of Fibonacci Sequence

Better algorithm? How to derive the general formula?

Properties of Fibonacci Sequence

Better algorithm? How to derive the general formula?

$$F_n = F_{n-1} + F_{n-2}$$

Observation: F_n is a linear combination of F_{n-1} and F_{n-2} . This hints us to use linear algebra to express the recurrence relation.

Properties of Fibonacci Sequence

Better algorithm? How to derive the general formula?

$$F_n = F_{n-1} + F_{n-2}$$

Observation: F_n is a linear combination of F_{n-1} and F_{n-2} . This hints us to use linear algebra to express the recurrence relation.

Proposition. Let $\{F_n\}$ be a Fibonacci sequence, then

$$\begin{pmatrix} F_{n+1} & F_n \\ F_n & F_{n-1} \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}^n$$

Proof by mathematical induction

Basis. $n = 1$:

$$\begin{pmatrix} F_2 & F_1 \\ F_1 & F_0 \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}$$

Proof (Induction Step)

Suppose for any n , the formula is correct, i.e.:

$$\begin{pmatrix} F_{n+1} & F_n \\ F_n & F_{n-1} \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}^n$$

Then for $n + 1$, according to the definition of Fibonacci sequence:

$$\begin{pmatrix} F_{n+2} & F_{n+1} \\ F_{n+1} & F_n \end{pmatrix} = \begin{pmatrix} F_{n+1} & F_n \\ F_n & F_{n-1} \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}$$

$$\text{Induction premise} \Rightarrow \left(\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}^n \right) \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} = \left(\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}^{n+1} \right)$$

Improved Algorithm via Fast Power

Let

$$M = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}$$

Compute M^n using generalized fast power algorithm

Time complexity

- The number of matrix multiplication $T(n) = \Theta(\log n)$
- Each matrix multiplication requires 8 number multiplication
- The overall complexity is $\Theta(\log n)$

Improved Algorithm via Fast Power

Let

$$M = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}$$

Compute M^n using generalized fast power algorithm

Time complexity

- The number of matrix multiplication $T(n) = \Theta(\log n)$
 - Each matrix multiplication requires 8 number multiplication
 - The overall complexity is $\Theta(\log n)$
-

Further improvement

- M can be diagonalized ($M = PM'P^{-1}$) \Rightarrow we could directly use fast power algorithm for better efficiency on basic computer step (matrix mul).

Outline

1 Fast Power/Exponentiation

2 Integer Multiplication

3 Matrix Multiplication

4 Polynomial Multiplication

Integer Addition

Addition. Given two n -bit integers a and b , compute $a + b$.

Subtraction. Given two n -bit integers a and b , compute $a - b$.

Grade-school algorithm. $\Theta(n)$ bit operations.

1	1	1	1	1	1	0	1
	1	1	0	1	0	1	0
+	0	1	1	1	1	1	0
	1	0	1	0	1	0	0

Remark. Grade-school addition and subtraction algorithms are asymptotically optimal.

Integer Multiplication

Multiplication. Given two n -bit integers a and b , compute $a \times b$.

Grade school method. $\Theta(n^2)$ bit operations

$\Theta(n^2)$ atomic bit multiplications + $\Theta(n^2)$ atomic bit additions

$$\begin{array}{r} & 1 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ \times & 0 & 1 & 1 & 1 & 1 & 1 & 0 & 1 \\ \hline & 1 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ & 1 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ & 1 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ & 1 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ & 1 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \end{array}$$

Divide-and-Conquer: First Attempt (1/2)

Divide Split two n -bit integer x and y into their left and right halves (low- and high-order bits). Let $m = n/2$.

x	x_L	x_R	$2^{n/2}x_L + x_R$
y	y_L	y_R	$2^{n/2}y_L + y_R$

Use bit shifting to compute

$$x_L = \lfloor x/2^m \rfloor, x_R = x \bmod 2^m$$
$$y_L = \lfloor y/2^m \rfloor, y_R = y \bmod 2^m$$

Example. $x = \underbrace{1011}_{x_L} \underbrace{0110}_{x_R} = 1011 \times 2^4 + 0110$

Divide-and-Conquer: First Attempt (2/2)

$$\begin{aligned} xy &= (2^{n/2}x_L + x_R)(2^{n/2}y_L + y_R) \\ &= 2^n \underline{x_L y_L} + 2^{n/2}(\underline{x_L y_R} + \underline{x_R y_L}) + \underline{x_R y_R} \end{aligned}$$

Conquer. Multiply four $n/2$ -bit integers, recursively. (significant operations)

Combine. Add and shift to obtain result.

$$T(n) = \underbrace{4T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n)}_{\text{add,shift}}$$

$$\text{master theorem (case 1)} \Rightarrow T(n) = \Theta(n^2)$$

(Subproblems) too many \sim Same running time as traditional grade school method, no progress in efficiency.

How can this method be sped up?

Gauss's Trick

Gauss once noticed that although the product of two complex numbers

$$(a + bi)(c + di) = \textcolor{red}{ac} - \textcolor{blue}{bd} + (bc + ad)i$$

seems involving 4 multiplications, it can in fact be done with 3:

$$bc + ad = (a + b)(c + d) - \textcolor{red}{ac} - \textcolor{blue}{bd}$$



Figure: Carl Friedrich Gauss

Karatsuba's Algorithm

In 1960, Kolmogorov conjectured grade-school multiplication algorithm is optimal in a seminar. Within a week, Karatsuba, then a 23-year-old student, found a much better algorithm thus disproving the conjecture. Kolmogorov was very excited about the discovery and published a paper in 1962.



Figure: Anatolii Karatsuba

Karatsuba algorithm: the first algorithm asymptotically faster than the quadratic “grade school” algorithm.

Reduce the Number of Subproblem

Idea. Exploit the dependence among subproblems via Gauss's trick

$$\underbrace{x_L y_R + x_R y_L}_{\text{middle term}} = (x_L + x_R)(y_L + y_R) - x_L y_L - x_R y_R$$

Algorithm 3: KARATSUBA(x, y, n)

- 1: **if** $n = 1$ **then return** $x \times y$;
 - 2: **else** $m \leftarrow \lceil n/2 \rceil$;
 - 3: $x_L \leftarrow \lfloor x/2^m \rfloor$; $x_R \leftarrow x \bmod 2^m$;
 - 4: $y_L \leftarrow \lfloor y/2^m \rfloor$; $y_R \leftarrow y \bmod 2^m$;
 - 5: $e \leftarrow \text{KARATSUBA}(x_L, y_L, m)$;
 - 6: $f \leftarrow \text{KARATSUBA}(x_R, y_R, m)$;
 - 7: $g \leftarrow \text{KARATSUBA}(x_L + x_R, y_L + y_R, m)$;
 - 8: **return** $2^{2m}e + 2^m(g - e - f) + f$;
-

Theory

Complexity Analysis. Now, the recurrence relation is

$$\left. \begin{array}{l} T(n) = 3T(n/2) + \Theta(n) \\ T(1) = 1 \end{array} \right\} \Rightarrow T(n) = \Theta(n^{\log_2 3}) = \Theta(n^{1.585})$$

Combining and dividing cost $f(n)$ is cheap $\leadsto h(n)$ dominates the overall complexity. The constant factor improvement from 4 to 3 occurs at the *every level of the recursion*, the compounding effect leads to a dramatically lower bound.

[Toom-Cook (1963)] faster generalization of Karatsuba's method

[Schönhage-Strassen (1971)] even faster, for sufficiently large n .

Some Notes

Practical note:

- It generally does not make sense to recurse all the way down to 1 bit. For most processors, 16- or 32-bit multiplication is single operation.
 - GNU Multiple Precision Library uses different algorithms depending on size of operands. (used in Maple, Mathematica, gcc, cryptography, . . .)
-

Theoretical recap (informal):

- Grad-school add/sub algorithms is optimal since they are already **local**
- Grad-school mul algorithm is not optimal since it is not very **local** (globally correlated): divide helps to shrink locality

Outline

1 Fast Power/Exponentiation

2 Integer Multiplication

3 Matrix Multiplication

4 Polynomial Multiplication

Inner Product

Inner product. Given two length n vectors $\mathbf{a} = (a_1, \dots, a_n)$ and $\mathbf{b} = (b_1, \dots, b_n)$, compute

$$c = \langle \mathbf{a}, \mathbf{b} \rangle = \sum_{i=1}^n a_i b_i$$

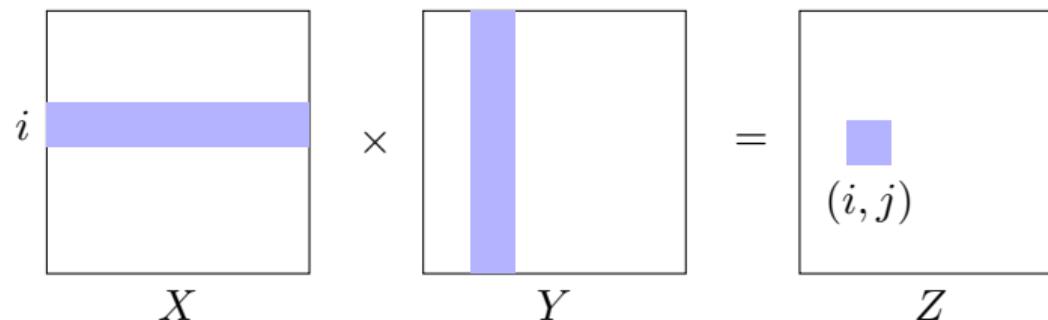
Grade school. $\Theta(n)$ arithmetic operations.

Remark. Grade-school dot product algorithm is asymptotically optimal.

Matrix Multiplication

Matrix multiplication. Given two $n \times n$ matrix X and Y , compute

$$Z = XY, Z_{ij} = \sum_{k=1}^n X_{ik} Y_{kj}$$



College-school method: $\Theta(n^3)$ arithmetic operations

- there are n^2 elements in Z
- computing each element requires n arithmetic multiplications

Is college-school matrix multiplication algorithm asymptotically optimal? Can divide-and-conquer strategy do better?

Naive Divide-and-Conquer

Strategy. Split matrix into blocks:

$$\begin{pmatrix} X_{11} & X_{12} \\ X_{21} & X_{22} \end{pmatrix} \begin{pmatrix} Y_{11} & Y_{12} \\ Y_{21} & Y_{22} \end{pmatrix} = \begin{pmatrix} Z_{11} & Z_{12} \\ Z_{21} & Z_{22} \end{pmatrix}$$

in which:

$$\begin{aligned} Z_{11} &= X_{11}Y_{11} + X_{12}Y_{21} & Z_{12} &= X_{11}Y_{12} + X_{12}Y_{22} \\ Z_{21} &= X_{21}Y_{11} + X_{22}Y_{21} & Z_{22} &= X_{21}Y_{12} + X_{22}Y_{22} \end{aligned}$$

Recurrence relation: master theorem (case 1)

$$T(n) = \underbrace{8T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n^2)}_{\text{add/form submatrices}} \quad \left. \begin{array}{l} \\ \\ T(1) = 1 \end{array} \right\} \Rightarrow T(n) = \Theta(n^3)$$

Breakthrough

College algorithm: $\Theta(n^3)$

Naive divide-and-conquer strategy: $\Theta(n^3)$ (unimpressive)

For a quite while, this was widely believed to the best running time possible, it was even proved that *in certain models* no algorithms can do better.

Great excitement: This efficiency can be further improved by some clever algebra.

Strassen Algorithm (1/3)

Volker Strassen first published this algorithm in 1969

- proved that the $\Theta(n^3)$ general matrix multiplication algorithm wasn't optimal
- faster than the standard matrix multiplication algorithm and is useful in practice for large matrices,
- inspired more research about matrix multiplication that led to faster approaches, e.g. the Coppersmith-Winograd algorithm.



Figure: Volker Strassen

Strassen Algorithm (2/3)

$$\begin{pmatrix} X_{11} & X_{12} \\ X_{21} & X_{22} \end{pmatrix} \begin{pmatrix} Y_{11} & Y_{12} \\ Y_{21} & Y_{22} \end{pmatrix} = \begin{pmatrix} Z_{11} & Z_{12} \\ Z_{21} & Z_{22} \end{pmatrix}$$

Define 7 instead matrix:

$$M_1 = X_{11}(Y_{12} - Y_{22})$$

$$M_2 = (X_{11} + X_{12})Y_{22}$$

$$M_3 = (X_{21} + X_{22})Y_{11}$$

$$M_4 = X_{22}(Y_{21} - Y_{11})$$

$$M_5 = (X_{11} + X_{22})(Y_{11} + Y_{22})$$

$$M_6 = (X_{12} - X_{22})(Y_{21} + Y_{22})$$

$$M_7 = (X_{11} - X_{21})(Y_{11} + Y_{12})$$

Express Z_{ij} via instead matrices

$$Z_{11} = M_5 + M_4 - M_2 + M_6$$

$$Z_{12} = M_1 + M_2$$

$$Z_{21} = M_3 + M_4$$

$$Z_{22} = M_5 + M_1 - M_3 - M_7$$

$$\begin{aligned} Z_{12} &= M_1 + M_2 = X_{11} \times (Y_{12} - Y_{22}) + (X_{11} + X_{12}) \times Y_{22} \\ &= X_{11} \times Y_{12} + X_{12} \times Y_{22} \end{aligned}$$

Strassen Algorithm (3/3)

Reduce the number of subproblems from 8 to 7

Recurrence relation for time complexity (18 is number of additions/subtraction performed at each application of the algorithm)

$$\left. \begin{array}{l} T(n) = 7T(n/2) + 18n^2 \\ T(1) = 1 \end{array} \right\} \Rightarrow T(n) = \Theta(n^{\log_2 7}) = \Theta(n^{2.8075})$$

Strassen Algorithm (3/3)

Reduce the number of subproblems from 8 to 7

Recurrence relation for time complexity (18 is number of additions/subtraction performed at each application of the algorithm)

$$\left. \begin{array}{l} T(n) = 7T(n/2) + 18n^2 \\ T(1) = 1 \end{array} \right\} \Rightarrow T(n) = \Theta(n^{\log_2 7}) = \Theta(n^{2.8075})$$

Q. What if n is not the power of 2?

A. Could pad matrices with zeros.

$$\begin{pmatrix} 1 & 2 & 3 & 0 \\ 4 & 5 & 7 & 0 \\ 7 & 8 & 9 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix} \begin{pmatrix} 10 & 11 & 12 & 0 \\ 13 & 14 & 15 & 0 \\ 16 & 17 & 18 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix} = \begin{pmatrix} 84 & 90 & 96 & 0 \\ 201 & 216 & 231 & 0 \\ 318 & 342 & 366 & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

More about Matrix Multiplication

The decomposition is *so tricky and intricate* that one wonders how Strassen was ever able to discover it!



More about Matrix Multiplication

The decomposition is *so tricky and intricate* that one wonders how Strassen was ever able to discover it!



Complexity of Matrix Multiplication

- Best upper-bound: $O(n^{2.376})$ — Coppersmith-Winograd algorithm
- Known lower-bound: $\Omega(n^2)$

More about Matrix Multiplication

The decomposition is *so tricky and intricate* that one wonders how Strassen was ever able to discover it!



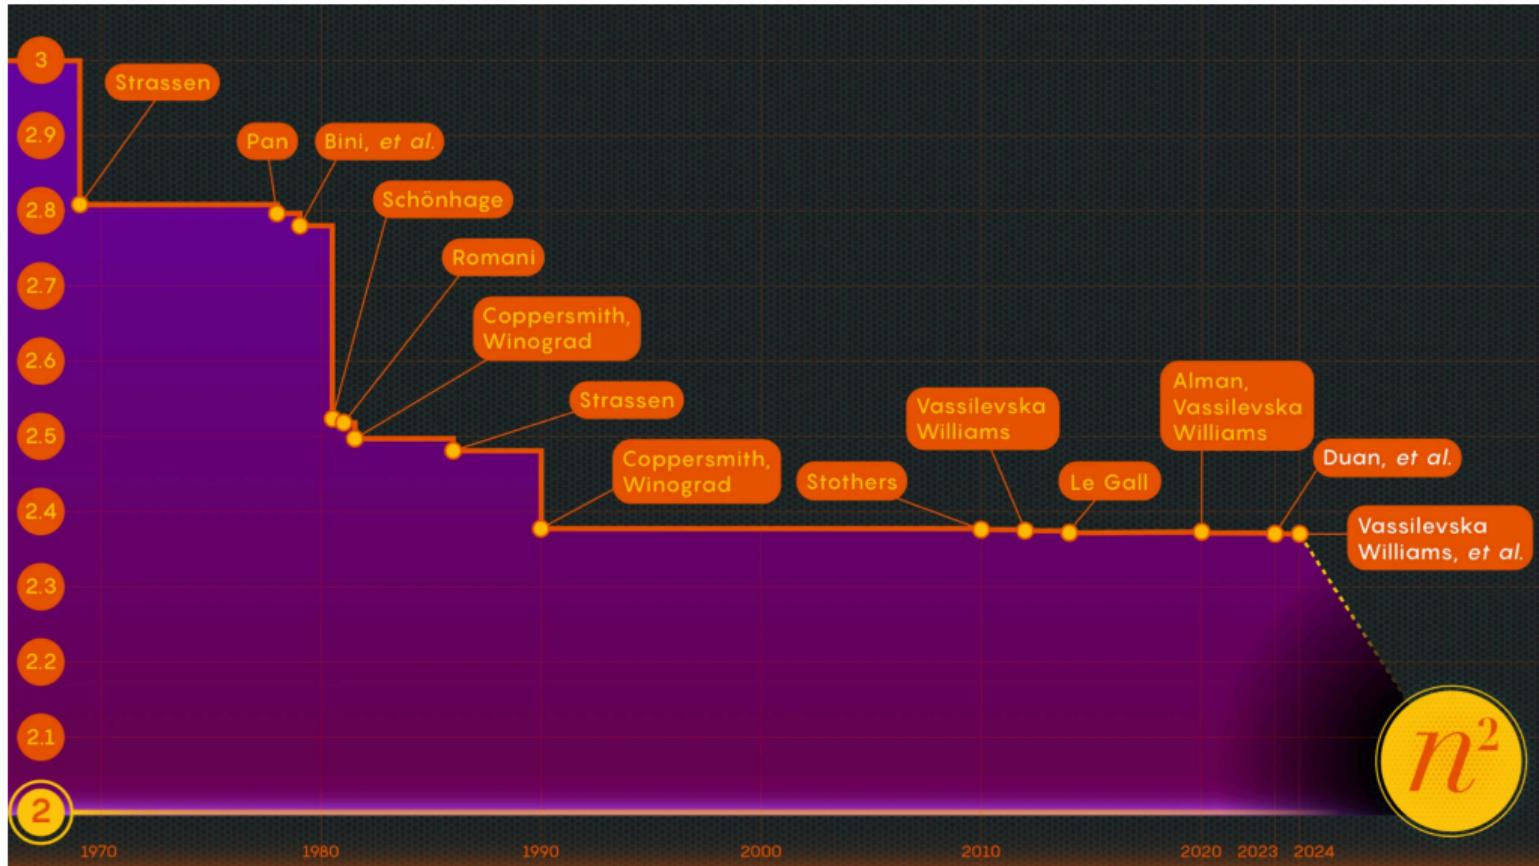
Complexity of Matrix Multiplication

- Best upper-bound: $O(n^{2.376})$ — Coppersmith-Winograd algorithm
- Known lower-bound: $\Omega(n^2)$

Applications

- scientific computation, image processing, data mining (regression, aggregation, decision tree)

Matrix Multiplication: SOTA



Matrix Multiplication: SOTA



Renfei Zhou

TCS Researcher

📍 Beijing, China

📍 Tsinghua University

About Me

I am a fourth-year undergraduate student studying theoretical computer science (TCS) in Yao Class, Tsinghua University. I am applying to Ph.D. programs this year.

I mainly work on (succinct) data structures and fast matrix multiplication. I also have broad interests in other fields of TCS, such as streaming algorithms, game theory, and online algorithms.

Selected Publications

Faster Matrix Multiplication via Asymmetric Hashing

Ran Duan *
Tsinghua University

Hongxun Wu †
UC Berkeley

Renfei Zhou ‡
Tsinghua University

New Bounds for Matrix Multiplication: from Alpha to Omega

Virginia Vassilevska Williams* Yinzhan Xu† Zixuan Xu‡ Renfei Zhou§

Matrix Multiplication: SOTA

矩阵乘法可能看起来是一个晦涩的问题，但它是一种基本的计算操作。它被融入了人们每天使用的大部分算法中，用于各种任务，从显示更清晰的计算机图形到解决网络理论中的物流问题。就像在计算的其他领域一样，速度至关重要。即使是微小的改进最终也可能大大减少所需要的时间、计算能力和金钱。[←](#)

段然团队修改了激光法标记块的方式，将 ω 的新上限设定在了 2.371866 左右，这要比 2020 年设定的上限 2.3728596 有所改进。这看起来是一个不大的变化，将上限降低了大约 0.001，但这是自 2010 年以来科学家们看到的最大进步。[←](#)

Le Gall 说：「要想进一步改进，就必须在 Coppersmith and Winograd 的原始方法基础上加以改进，而这种方法自 1987 年以来就没有真正改变过。」但到目前为止，还没有人提出更好的方法。也许根本就没有。[←](#)

周任飞说：「改进 ω 实际上是理解这个问题的一部分。如果我们能很好地理解这个问题，就能设计出更好的算法。不过，人们对这个古老问题的理解还处于非常初级的阶段。」[←](#)

Outline

- 1 Fast Power/Exponentiation
- 2 Integer Multiplication
- 3 Matrix Multiplication
- 4 Polynomial Multiplication

Motivation

We have studied how to multiply

- Integers: Gauss's trick
- Matrix: Strassen algorithm

Motivation

We have studied how to multiply

- Integers: Gauss's trick
- Matrix: Strassen algorithm

How to multiply polynomials?

Motivation

We have studied how to multiply

- Integers: Gauss's trick
- Matrix: Strassen algorithm

How to multiply polynomials?

Applications of polynomial multiplication

- Fastest polynomial multiplication implies fastest integers multiplication
 - polynomials and binary integers are quite similar — just replace the variable x by the base 2 and watch out for carries
- Multiplying polynomials is crucial for *signal processing*

Polynomials: Coefficient Representation

Polynomial. [coefficient representation]

$$A(x) = a_0 + a_1x + a_2x^2 + \cdots + a_{n-1}x^{n-1}$$

$$B(x) = b_0 + b_1x + b_2x^2 + \cdots + b_{n-1}x^{n-1}$$

Polynomial Operation

Add. $\Theta(n)$

$$A(x) + B(x) = (a_0 + b_0) + (a_1 + b_1)x + \cdots + (a_{n-1} + b_{n-1})x^{n-1}$$

Polynomial Operation

Add. $\Theta(n)$

$$A(x) + B(x) = (a_0 + b_0) + (a_1 + b_1)x + \cdots + (a_{n-1} + b_{n-1})x^{n-1}$$

Evaluate. three choices:

- Naive algorithm. compute each term one by one: $\Theta(n^2)$
- Caching algorithm. cache x^i : $\Theta(n)$
- Horner algorithm. $a_0 + (x(a_1 + x(a_2 + \cdots + x(a_{n-2} + x(a_{n-1}))))))$: $\Theta(n)$
秦九韶 discovered this algorithm hundreds of years earlier
参考链接: <https://zhuanlan.zhihu.com/p/22166332>

Polynomial Operation

Add. $\Theta(n)$

$$A(x) + B(x) = (a_0 + b_0) + (a_1 + b_1)x + \cdots + (a_{n-1} + b_{n-1})x^{n-1}$$

Evaluate. three choices:

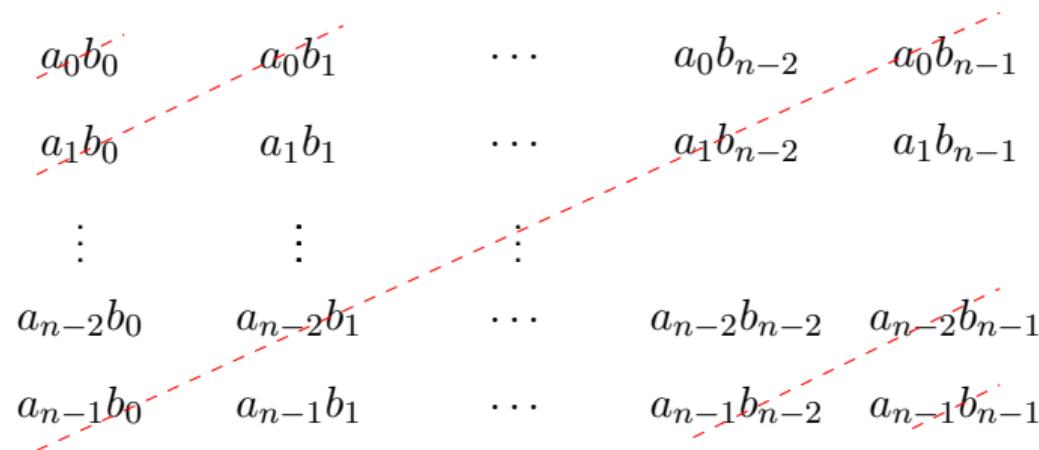
- Naive algorithm. compute each term one by one: $\Theta(n^2)$
- Caching algorithm. cache x^i : $\Theta(n)$
- Horner algorithm. $a_0 + (x(a_1 + x(a_2 + \cdots + x(a_{n-2} + x(a_{n-1}))))))$: $\Theta(n)$
秦九韶 discovered this algorithm hundreds of years earlier
参考链接: <https://zhuanlan.zhihu.com/p/22166332>

Multiply (convolve). $\Theta(n^2)$ using brute force algorithm

$$\begin{aligned} A(x) \times B(x) &= a_0b_0 + (a_0b_1 + a_1b_0)x + \cdots + a_{n-1}b_{n-1}x^{2n-2} \\ &= \sum_{i=0}^{2n-2} c_i x^i, \text{ where } c_i = \sum_{j=0}^i a_j b_{i-j} \end{aligned}$$

Pictorial View of Convolution

$$(c_0, \dots, c_{2n-2}) = (a_0, \dots, a_{n-1}) \circledast (b_0, \dots, b_{n-1})$$

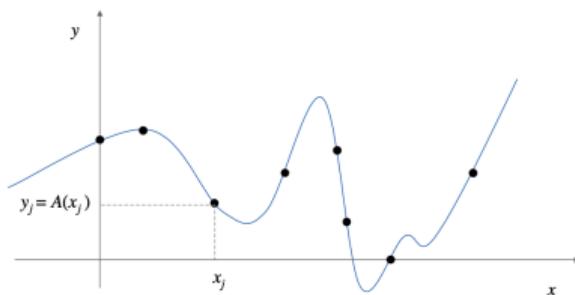


Polynomials: Point-Value Representation

Fundamental Theorem of Algebra [Gauss, PhD thesis]

Every non-zero, single-variable, degree n polynomial with complex coefficients has exactly n complex roots.

Corollary. A degree $n - 1$ polynomial $A(x)$ is uniquely specified by its evaluations at n distinct points.



Suppose another degree $n - 1$ polynomial $A'(x)$ has same evaluations at n distinct points as $A(x)$
 $\Rightarrow A(x) - A'(x)$ is of at most degree $n - 1$ but has n roots \leadsto contradicts to the fundamental theorem of algebra

Polynomials: Point-Value Representation

Polynomial. [point-value representation]

$$A(x) : (x_0, y_0), \dots, (x_{n-1}, y_{n-1})$$

$$B(x) : (x_0, z_0), \dots, (x_{n-1}, z_{n-1})$$

Polynomial Operation

Add. $\Theta(n)$ add operations.

$$A(x) + B(x) : (x_0, y_0 + z_0), \dots, (x_{n-1}, y_{n-1} + z_{n-1})$$

Multiply (convolve). $\Theta(n)$, but need $2n - 1$ points.

$$A(x) \times B(x) : (x_0, y_0 \times z_0), \dots, (x_{2n-2}, y_{2n-2} \times z_{2n-2})$$

Evaluate. $\Theta(n^2)$ using Lagrange's formula

$$A(x) = \sum_{k=0}^{n-1} y_k \frac{\prod_{j \neq k} (x - x_j)}{\prod_{j \neq k} (x_k - x_j)}$$

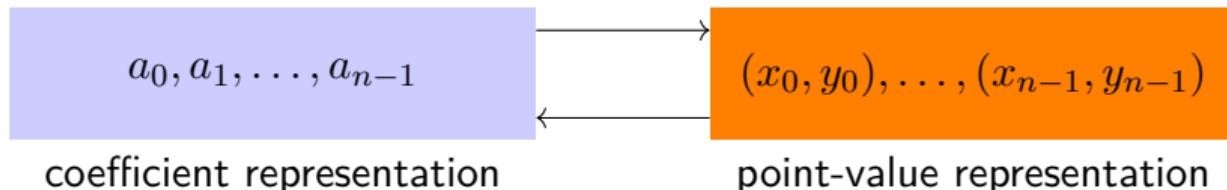
Conversion Between Two Representations

Trade-off. Fast evaluation or Fast multiplication

representation	multiply	evaluate
coefficient	$\Theta(n^2)$	$\Theta(n)$
point-value	$\Theta(n)$	$\Theta(n^2)$



Goal. Efficient conversion between two representations \Rightarrow enjoy the best of two worlds: all operations are fast!



Conversion Between Two Representations: Evaluation

Coefficient \Rightarrow Point-value

Given $A(x) = a_0 + a_1x + \cdots + a_{n-1}x^{n-1}$, evaluate it at n distinct points x_0, \dots, x_{n-1} .

$$\begin{bmatrix} y_0 \\ y_1 \\ y_2 \\ \vdots \\ y_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \cdots & x_1^{n-1} \\ 1 & x_2 & x_2^2 & \cdots & x_2^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n-1} & x_{n-1}^2 & \cdots & x_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Running time. $\Theta(n^2)$ for matrix-vector multiply (or n Horner's).

Conversion Between Two Representations: Interpolation

Point-value \Rightarrow Coefficient

Given n distinct points x_0, \dots, x_{n-1} and values (y_0, \dots, y_{n-1}) , find unique polynomial $A(x) = a_0 + a_1x + \dots + a_{n-1}x^{n-1}$, that has given values at given points.

$$\begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & x_0 & x_0^2 & \dots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \dots & x_1^{n-1} \\ 1 & x_2 & x_2^2 & \dots & x_2^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n-1} & x_{n-1}^2 & \dots & x_{n-1}^{n-1} \end{bmatrix}^{-1} \begin{bmatrix} y_0 \\ y_1 \\ y_2 \\ \vdots \\ y_{n-1} \end{bmatrix}$$

Vandermonde matrix is invertible iff x_i 's are distinct.

Running time. $\Theta(n^3)$ for Gaussian elimination

Restate Our Goal

Both known conversions are inefficient

- coefficients \Rightarrow point-value: $\Theta(n^2)$
- point-value \Rightarrow coefficients: $\Theta(n^3)$

More efficient conversions are needed.

Restate Our Goal

Both known conversions are inefficient

- coefficients \Rightarrow point-value: $\Theta(n^2)$
- point-value \Rightarrow coefficients: $\Theta(n^3)$

More efficient conversions are needed.

Next, we begin with the first direction. We restate our goal:

Given n coefficients, computing n point-value tuples quickly.

Restate Our Goal

Both known conversions are inefficient

- coefficients \Rightarrow point-value: $\Theta(n^2)$
- point-value \Rightarrow coefficients: $\Theta(n^3)$

More efficient conversions are needed.

Next, we begin with the first direction. We restate our goal:

Given n coefficients, computing n point-value tuples quickly.

The optimal complexity of polynomial evaluation algorithm is $\Theta(n)$. Thus, it seems that $\Theta(n^2)$ complexity for the above goal is inevitable.

Restate Our Goal

Both known conversions are inefficient

- coefficients \Rightarrow point-value: $\Theta(n^2)$
- point-value \Rightarrow coefficients: $\Theta(n^3)$

More efficient conversions are needed.

Next, we begin with the first direction. We restate our goal:

Given n coefficients, computing n point-value tuples quickly.

The optimal complexity of polynomial evaluation algorithm is $\Theta(n)$. Thus, it seems that $\Theta(n^2)$ complexity for the above goal is inevitable.

High-level key idea: impose **structure** to enhance locality \leadsto reduce complexity

Divide-and-Conquer for Evaluation

$$A(x) = a_0 + a_1x + a_2x^2 + a_3x^3 + a_4x^4 + a_5x^5 + a_6x^6 + a_7x^7$$

two choices for dividing: frequency vs. time

Decimation in frequency. Break polynomial into low and high powers.

$$A_{\text{low}}(x) = a_0 + a_1x + a_2x^2 + a_3x^3$$

$$A_{\text{high}}(x) = a_4 + a_5x + a_6x^2 + a_7x^3$$

$$A(x) = A_{\text{low}}(x) + x^4 A_{\text{high}}(x).$$

Decimation in time. Break polynomial into even and odd powers.

$$A_{\text{even}}(x) = a_0 + a_2x + a_4x^2 + a_6x^3$$

$$A_{\text{odd}}(x) = a_1 + a_3x + a_5x^2 + a_7x^3$$

$$A(x) = A_{\text{even}}(x^2) + x A_{\text{odd}}(x^2)$$

radix-2 decimation-in-time (DIT)

Clarification

We emphasize that the goal of dividing is not to improve the efficiency of polynomial evaluation at a single point, cause it is already optimal.

The ultimate goal is to improve the efficiency of evaluation of n points as a whole task.

First Attempt

Naive idea. Randomly pick n distinct points x_0, \dots, x_{n-1} , then compute $A(x)$ via $A_{\text{even}}(x^2) + xA_{\text{odd}}(x^2)$.

- $T(n)$: evaluate degree- $(n - 1)$ polynomial at n points
- $E(n)$: evaluate degree- $(n - 1)$ polynomial at 1 point

First Attempt

Naive idea. Randomly pick n distinct points x_0, \dots, x_{n-1} , then compute $A(x)$ via $A_{\text{even}}(x^2) + xA_{\text{odd}}(x^2)$.

- $T(n)$: evaluate degree- $(n - 1)$ polynomial at n points
- $E(n)$: evaluate degree- $(n - 1)$ polynomial at 1 point

Issue. Efficiency does not improve

- Evaluating $A(x)$ of degree- $(n - 1)$ at n points: $T(n) = n \cdot E(n)$
- Evaluating $A_{\text{even}}(x)$ and $A_{\text{odd}}(x)$ both of degree- $(n/2 - 1)$ at n points:
 $2 \times n \cdot E(n/2) = 2n \cdot E(n/2)$

$E(n)$ is a linear function \leadsto no efficiency improvement

- The root is that the size of problem does not fully shrink by half.

Solution. Reduce the number of evaluated points

Basic Idea (1/2)

Basic idea. Introduce simple structure by choosing the n points to be positive-negative pairs, that is,

$$\pm x_0, \pm x_1, \dots, \pm x_{n/2-1}$$

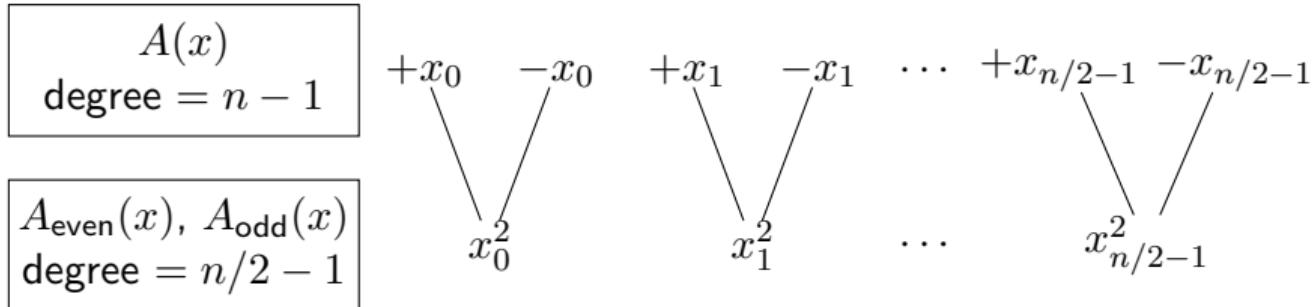
Note that the **even powers** of x_i coincide with those of $-x_i \Rightarrow$ the computations required for each $A(x_i)$ and $A(-x_i)$ overlap a lot.

$$A(x_i) = A_{\text{even}}(x_i^2) + x_i A_{\text{odd}}(x_i^2)$$

$$A(-x_i) = A_{\text{even}}(x_i^2) - x_i A_{\text{odd}}(x_i^2)$$

Now, evaluating degree- $(n - 1)$ polynomial $A(x)$ at n paired points $\pm x_0, \dots, \pm x_{n/2-1} \Rightarrow$ evaluating degree- $(n/2 - 1)$ polynomials $A_{\text{even}}(x)$ and $A_{\text{odd}}(x)$ at just **$n/2$** points: $x_0^2, \dots, x_{n/2-1}^2$.

Basic Idea (2/2)



Now, the original problem of size n is in this way recast as two subproblems of size $n/2$ followed by some linear-time arithmetic.

$T(n)$: evaluate a degree $(n - 1)$ polynomial at n points

- If we could recurse, we would get a divide-and-conquer procedure with running time:

$$T(n) = 2T(n/2) + \Theta(n)$$

which is $\Theta(n \log n)$, exactly what we want.

Technical Hurdle

Technical hurdle. The plus-minus trick only works at the top level of the recursion.

- To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ *themselves* to be plus-minus pairs.

Technical Hurdle

Technical hurdle. The plus-minus trick only works at the top level of the recursion.

- To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ *themselves* to be plus-minus pairs.

But how can a square to be negative?

Technical Hurdle

Technical hurdle. The plus-minus trick only works at the top level of the recursion.

- To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ themselves to be plus-minus pairs.

But how can a square to be negative?



Technical Hurdle

Technical hurdle. The plus-minus trick only works at the top level of the recursion.

- To recurse at the next level, we need the $n/2$ evaluation points $x_0^2, x_1^2, \dots, x_{n/2-1}^2$ themselves to be plus-minus pairs.

But how can a square to be negative?

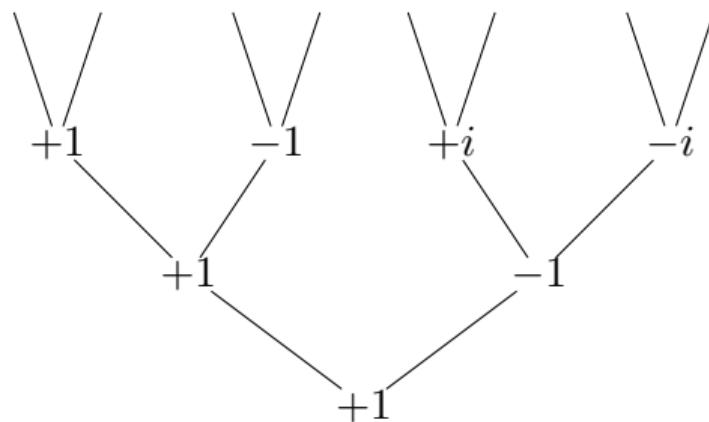


Unless, of course, we use complex numbers.

Which Complex Numbers?

Fine, but which complex numbers? Let us figure out by “reverse engineer” the process.

- At the bottom of the recursion, we have a single point, say, 1.
- In the level above it must consists of its square roots, ± 1 .
- The next level up is $(+1, -1)$ and $(+i, -i)$, until we reach $n = 2^k$ leaf nodes.



The Choice of n Complex Numbers

An n -th root of unity is a complex number such that $x^n = 1$.

The Choice of n Complex Numbers

An n -th root of unity is a complex number such that $x^n = 1$.

Fact. The n -th roots of unity are $\omega^0 = 1, \omega^1, \omega^2, \dots, \omega^{n-1}$, where
 $\omega = e^{2\pi i/n} = \cos \frac{2\pi}{n} + i \sin \frac{2\pi}{n}$

The Choice of n Complex Numbers

An n -th root of unity is a complex number such that $x^n = 1$.

Fact. The n -th roots of unity are $\omega^0 = 1, \omega^1, \omega^2, \dots, \omega^{n-1}$, where
 $\omega = e^{2\pi i/n} = \cos \frac{2\pi}{n} + i \sin \frac{2\pi}{n}$

Proof. $(\omega^k)^n = (e^{2\pi ik/n})^n = (e^{\pi i})^{2k} = (-1)^{2k} = 1$

$$e^{ix} = \cos x + i \sin x$$

If n is even, the n -th roots are plus-minus paired, $\omega^{n/2+j} = -\omega^j$

- Squaring them produces the $(n/2)$ -th roots of unity: $v^0, v^1, \dots, v^{n/2-1}$, where
 $v = \omega^2 = e^{4\pi i/n}$.

The Choice of n Complex Numbers

An n -th root of unity is a complex number such that $x^n = 1$.

Fact. The n -th roots of unity are $\omega^0 = 1, \omega^1, \omega^2, \dots, \omega^{n-1}$, where
 $\omega = e^{2\pi i/n} = \cos \frac{2\pi}{n} + i \sin \frac{2\pi}{n}$

Proof. $(\omega^k)^n = (e^{2\pi ik/n})^n = (e^{\pi i})^{2k} = (-1)^{2k} = 1$

$$e^{ix} = \cos x + i \sin x$$

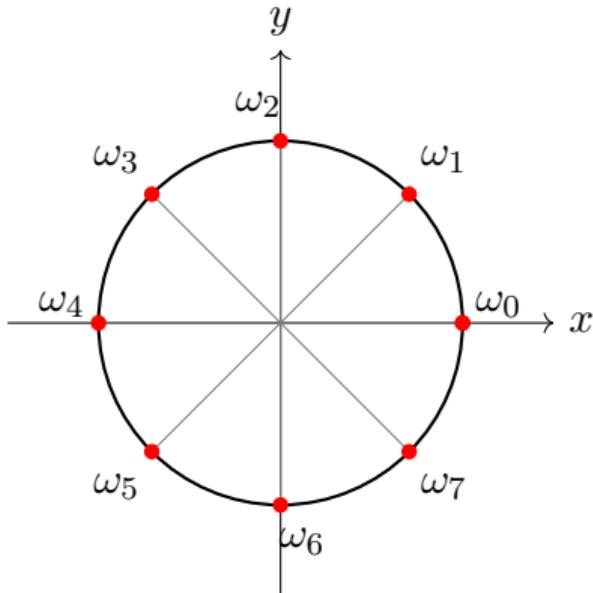
If n is even, the n -th roots are plus-minus paired, $\omega^{n/2+j} = -\omega^j$

- Squaring them produces the $(n/2)$ -th roots of unity: $v^0, v^1, \dots, v^{n/2-1}$, where
 $v = \omega^2 = e^{4\pi i/n}$.

If we start with $\omega^0, \omega^1, \omega^2, \dots, \omega^{n-1}$ for some $n = 2^k$, then at k -level of recursion we will have the $(n/2^k)$ -th roots of unity.

- All these sets of roots are plus-minus paired \Rightarrow Divide-and-conquer algorithm will work perfectly

Demo of $n = 8$



$$\omega_0 = 1$$

$$\omega_1 = e^{\frac{\pi}{4}i} = \frac{\sqrt{2}}{2} + \frac{\sqrt{2}}{2} \cdot i$$

$$\omega_2 = e^{\frac{\pi}{2}i} = i$$

$$\omega_3 = e^{\frac{3\pi}{4}i} = -\frac{\sqrt{2}}{2} + \frac{\sqrt{2}}{2} \cdot i$$

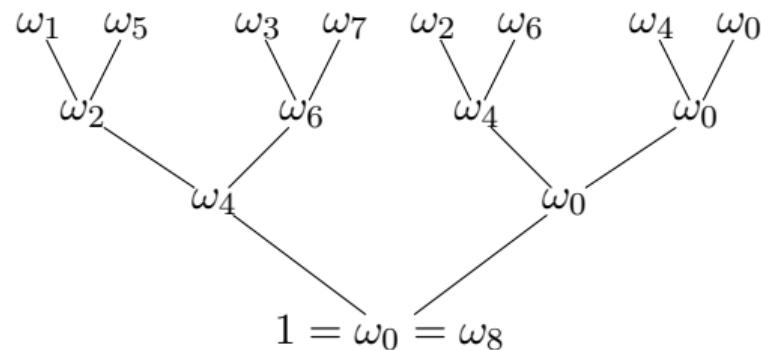
$$\omega_4 = e^{\pi i} = -1$$

$$\omega_5 = e^{\frac{5\pi}{4}i} = -\frac{\sqrt{2}}{2} - \frac{\sqrt{2}}{2} \cdot i$$

$$\omega_6 = e^{\frac{3\pi}{2}i} = -i$$

$$\omega_7 = e^{\frac{7\pi}{4}i} = \frac{\sqrt{2}}{2} - \frac{\sqrt{2}}{2} \cdot i$$

Recursion Structure and FFT



DFT: Fourier Matrix $M_n(\omega)$

$$\begin{bmatrix} y_0 \\ y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & \omega_0 & \omega_0^2 & \dots & \omega_0^{n-1} \\ 1 & \omega_1 & \omega_1^2 & \dots & \omega_1^{n-1} \\ 1 & \omega_2 & \omega_2^2 & \dots & \omega_2^{n-1} \\ 1 & \omega_3 & \omega_3^2 & \dots & \omega_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega_{n-1} & \omega_{n-1}^2 & \dots & \omega_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

Fast Fourier Transform (FFT)

Refined Goal. Evaluate $A(x) = a_0 + \cdots + a_{n-1}x^{n-1}$ at its n -th root of unity:
 $\omega^0, \omega^1, \dots, \omega^{n-1}$

Divide. Break up polynomial into even and odd powers:

$$A_{\text{even}}(x) = a_0 + a_2x + a_4x^2 + \cdots + a_{n-2}x^{n/2-1}$$

$$A_{\text{odd}}(x) = a_1 + a_3x + a_5x^2 + \cdots + a_{n-1}x^{n/2-1}$$

$$A(x) = A_{\text{even}}(x^2) + xA_{\text{odd}}(x^2)$$

Conquer. Evaluate $A_{\text{even}}(x)$ and $A_{\text{odd}}(x)$ at the $n/2$ -th roots of unity:

$$v^0, v^1, \dots, v^{n/2-1}$$

$$v^k = (\omega^k)^2$$

Combine.

$$A(\omega^k) = A_{\text{even}}(v^k) + \omega^k A_{\text{odd}}(v^k), 1 \leq k < n/2$$

$$A(\omega^{k+n/2}) = A_{\text{even}}(v^k) - \omega^k A_{\text{odd}}(v^k), 1 \leq k < n/2$$

$$v^k = (\omega^{k+n/2})^2 \quad (\omega^{k+n/2}) = -\omega^k$$

Pseudocode of FFT Algorithm

Algorithm 4: FFT(A, n, ω)

Input: coefficient representation of degree $n - 1$ polynomial A , principal n -th root of unity $\omega = e^{2\pi i/n}$

Output: value representation $A(\omega^0), \dots, A(\omega^{n-1})$

```
1: if  $n = 1$  then return  $a_0$ ;  
2: express  $A(x) = A_{\text{even}}(x^2) + xA_{\text{odd}}(x^2)$ ;  
3: FFT( $A_{\text{even}}, \frac{n}{2}, \omega^2$ )  $\rightarrow (A_{\text{even}}((\omega^2)^0), \dots, A_{\text{even}}((\omega^2)^{n/2-1}))$ ;  
4: FFT( $A_{\text{odd}}, \frac{n}{2}, \omega^2$ )  $\rightarrow (A_{\text{odd}}((\omega^2)^0), \dots, A_{\text{odd}}((\omega^2)^{n/2-1}))$ ;  
5: for  $j = 0$  to  $n - 1$  do  
6:    $A(\omega^j) = A_{\text{even}}(\omega^{2j}) + \omega^j A_{\text{odd}}(\omega^{2j})$  //  $\Theta(n)$ ;  
7: end  
8: return  $A(\omega^0), \dots, A(\omega^{n-1})$ ;
```

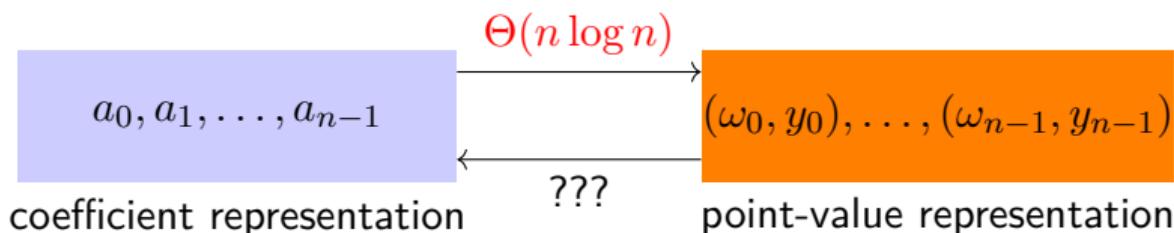
FFT Summary

Theorem. Assume $n = 2^k$. FFT algorithm evaluates a degree $n - 1$ polynomial at each of the n -th roots of unity in $\Theta(n \log n)$ steps.

Running time

$$T(n) = 2T(n/2) + \Theta(n) \Rightarrow T(n) = \Theta(n \log n)$$

Essence: choose n points with special structure to accelerate DFT computation.



Recap

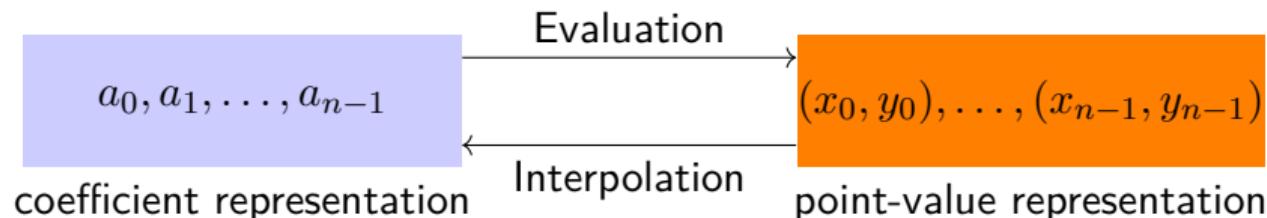
We first developed a high-level method for multiplying polynomials

coefficient representation \Rightarrow point-value representation

Point-value representation makes it trivial to multiply polynomials, but the input-output form of algorithm is specified as coefficient representation.

- So we designed FFT: coefficient \Rightarrow point-value in time just $\Theta(n \log n)$, where the points $\{x_i\}_n$ are complex n -th roots of unity $(1, \omega, \omega^2, \dots, \omega^{n-1})$.

$$\langle \text{values} \rangle = \text{FFT}(\langle \text{coefficients} \rangle, \omega)$$



Interpolation

The last remaining piece of the puzzle is the inverse operation — interpolation. It turns out amazingly that:

$$\langle \text{coefficients} \rangle = \frac{1}{n} \text{FFT}(\langle \text{values} \rangle, \omega^{-1})$$

Interpolation is thus solved in the most simple and elegant way, using the same FFT algorithm, but called with ω^{-1} in place of ω !

This might seem like a miraculous coincidence, but it will make a lot more sense when recasting polynomial operations in the language of linear algebra.

Inverse Discrete Fourier Transform

Point-value \Rightarrow Coefficient

Given n distinct points x_0, \dots, x_{n-1} and values y_0, \dots, y_{n-1} , find unique polynomial $A(x) = a_0 + a_1x + \dots + a_{n-1}x^{n-1}$, that has given values at given points.

Inverse DFT: Fourier Matrix inverse $F_n(\omega)^{-1}$

$$\begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \\ \vdots \\ a_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & \omega_0 & \omega_0^1 & \dots & \omega_0^{n-1} \\ 1 & \omega_1 & \omega_1^2 & \dots & \omega_1^{n-1} \\ 1 & \omega_2 & \omega_2^2 & \dots & \omega_2^{n-1} \\ 1 & \omega_3 & \omega_3^2 & \dots & \omega_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega_{n-1} & \omega_{n-1}^2 & \dots & \omega_{n-1}^{n-1} \end{bmatrix}^{-1} \begin{bmatrix} y_0 \\ y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_{n-1} \end{bmatrix}$$

$$F_n(\omega) = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega^1 & \omega^2 & \dots & \omega^{n-1} \\ 1 & \omega^2 & \omega^4 & \dots & \omega^{2(n-1)} \\ 1 & \omega^3 & \omega^6 & \dots & \omega^{3(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-1} & \omega^{2(n-1)} & \dots & \omega^{(n-1)(n-1)} \end{bmatrix}$$

$$F_n(\omega^{-1}) = \begin{bmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & \omega^{-1} & \omega^{-2} & \dots & \omega^{-(n-1)} \\ 1 & \omega^{-2} & \omega^{-4} & \dots & \omega^{-2(n-1)} \\ 1 & \omega^{-3} & \omega^{-6} & \dots & \omega^{-3(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{-(n-1)} & \omega^{-2(n-1)} & \dots & \omega^{-(n-1)(n-1)} \end{bmatrix}$$

Key Fact

$$G_n(\omega) := \frac{1}{n} F_n(\omega^{-1}) = F_n(\omega)^{-1}$$

Claim. F_n and G_n are inverses

Proof. Examine $F_n G_n$

$$(F_n G_n)_{kk'} = \frac{1}{n} \sum_{j=0}^{n-1} \omega^{kj} \omega^{-jk'} = \frac{1}{n} \sum_{j=0}^{n-1} \omega^{(k-k')j} = \begin{cases} 1 & \text{if } k = k' \\ 0 & \text{otherwise} \end{cases}$$

Summation lemma. Let ω be the principal n -th root of unity. Then

$$\sum_{j=0}^{n-1} \omega^{kj} = \begin{cases} n & \text{if } k = 0 \bmod n \\ 0 & \text{otherwise} \end{cases}$$

- If k is the multiple of n then $\omega^k = 1 \Rightarrow$ series sums to n
- Each ω^k is a root of $x^n - 1$

$$x^n - 1 = (x - 1)(1 + x + x^2 + \cdots + x^{n-1}) \Rightarrow \text{if } \omega^k \neq 1 \text{ we have:}$$
$$1 + \omega^k + \omega^{k(2)} + \cdots + \omega^{k(n-1)} = 0 \Rightarrow \text{series sums to 0}$$

Consequence

To compute inverse FFT, apply same algorithm but use

- $\omega^{-1} = e^{-2\pi i/n}$ as principal n -th root of unity (and divide the result by n).
- switch the role of $\langle a_0, \dots, a_{n-1} \rangle$ and $\langle y_0, \dots, y_{n-1} \rangle$

Interpolation Resolved

$$\langle \text{values} \rangle = \text{FFT}(\langle \text{coefficients} \rangle, \omega)$$

$$\begin{bmatrix} A(\omega_0) \\ A(\omega_1) \\ A(\omega_2) \\ \vdots \\ A(\omega_{n-1}) \end{bmatrix} = \begin{bmatrix} 1 & \omega_0 & \omega_0^2 & \dots & \omega_0^{n-1} \\ 1 & \omega_1 & \omega_1^2 & \dots & \omega_1^{n-1} \\ 1 & \omega_2 & \omega_2^2 & \dots & \omega_2^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega_{n-1} & \omega_{n-1}^2 & \dots & \omega_{n-1}^{n-1} \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_{n-1} \end{bmatrix}$$

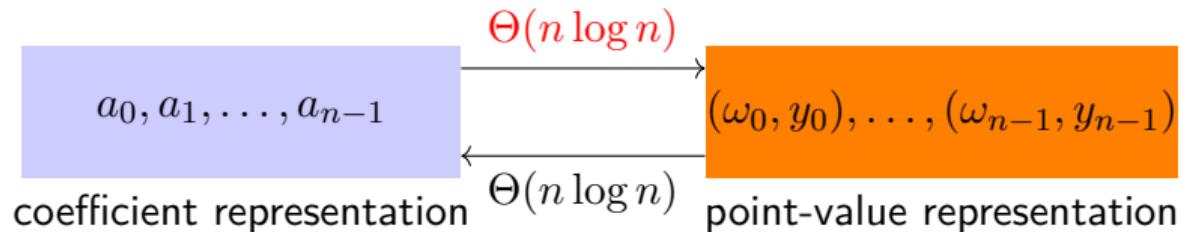
$$\langle \text{coefficients} \rangle = \frac{1}{n} \text{FFT}(\langle \text{values} \rangle, \omega^{-1})$$

$$\begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_{n-1} \end{bmatrix} = \frac{1}{n} \begin{bmatrix} 1 & (\omega_0^{-1})^1 & (\omega_0^{-1})^2 & \dots & (\omega_0^{-1})^{n-1} \\ 1 & (\omega_1^{-1})^1 & (\omega_1^{-1})^2 & \dots & (\omega_1^{-1})^{n-1} \\ 1 & (\omega_2^{-1})^1 & (\omega_2^{-1})^2 & \dots & (\omega_2^{-1})^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & (\omega_{n-1}^{-1})^1 & (\omega_{n-1}^{-1})^2 & \dots & (\omega_{n-1}^{-1})^{n-1} \end{bmatrix} \begin{bmatrix} A(\omega_0) \\ A(\omega_1) \\ A(\omega_2) \\ \vdots \\ A(\omega_{n-1}) \end{bmatrix}$$

Inverse FFT Summary

Theorem. Assume $n = 2^k$. Inverse FFT algorithm interpolated a degree $n - 1$ polynomial given values at each of the n -th roots of unity in $\Theta(n \log n)$ steps.

Running time. Almost the same algorithm as FFT.



Polynomial Multiplication

Theorem. Two degree $(n - 1)$ -polynomials can be multiplied in $\Theta(n \log n)$ steps.
(pad with 0s to make n a power of 2)

coefficient representation

$$\begin{aligned} a_0, a_1, \dots, a_{n-1} \\ b_0, b_1, \dots, b_{n-1} \end{aligned}$$

coefficient representation

$$c_0, c_1, \dots, c_{2n-2}, c_{2n-1} = 0$$

2 FFTs

$\Theta(n \log n)$

1 inverse FFT

$\Theta(n \log n)$

$$\begin{aligned} A(\omega_0), \dots, A(\omega_{2n-1}) \\ B(\omega_0), \dots, B(\omega_{2n-1}) \end{aligned}$$

$$\frac{\text{point-value}}{\text{multiplication}} \xrightarrow{\Theta(n)} \frac{\Theta(n)}{C(\omega_0), \dots, C(\omega_{2n-1})}$$

Actually, $2n - 1$ point-value tuples are sufficient.
But, FFT requires the input size to be 2^k , so is the output size

Remarks

Standard FFT. Evaluating degree $(n - 1)$ - $A(x)$ at its n -th root of unity $\omega^0, \omega^1, \dots, \omega^{n-1}$ by evaluating degree $n/2 - 1$ polynomials $A_{\text{even}}(x)$ and $A_{\text{odd}}(x)$ at their $n/2$ -th root of unity.

We choose the degree of polynomial as input size, since it determines the depth of recursion call.

Standard FFT can be easily extended to evaluating degree $(n - 1)$ polynomial $A(x)$ at its $2n$ -th root of unity $\omega^0, \omega^1, \dots, \omega^{2n-1}$ by evaluating degree $(n/2 - 1)$ polynomials $A_{\text{even}}(x)$ and $A_{\text{odd}}(x)$ at their n -th root of unity.

We still choose the degree of polynomial as input size, the recurrence relation is similar,

$$f(n) : \Theta(n) \sim \Theta(2n)$$

The overall complexity does not change in asymptotic sense.

Extension of FFT

FFT works in the field of complex numbers \mathbb{C} , the roots might be complex numbers \rightsquigarrow precision loss is inevitable

Sometimes we only need to work in a finite field, e.g. $\mathbb{F} = \mathbb{Z}/p$, the integers modulo a prime p .

- Primitive n -th roots of unity exist whenever n divides $p - 1$, so we have $p = \xi n + 1$ for a positive integer ξ .
- Specially, let ω be a primitive $(p - 1)$ -th root of unity, then an n -th root of unity — α can be found by letting $\alpha = \omega^\xi$

Extension \Rightarrow Number-Theoretic Transform (NTT): obtained by specializing the discrete Fourier transform to \mathbb{F} .

- no precision loss & much faster
- used extensively in the implementation of SNARK, e.g. `libsrank`

`libfqfft`

C++ library for FFTs in Finite Fields

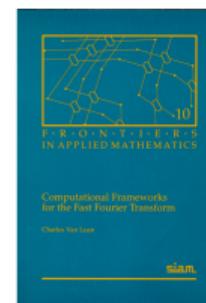
`libfqfft` is a C++ library for Fast Fourier Transforms (FFTs) in finite fields with multithreading support (via OpenMP). The library is developed by [SCIPR Lab](#) and contributors (see [AUTHORS](#) file) and is released under the MIT License (see [LICENSE](#) file). The library provides functionality for fast multipoint polynomial evaluation, fast polynomial interpolation, and fast computation of Lagrange polynomials. Check out the [Performance](#) section for memory, runtime, and field operation profiling data.

Applications of FFT

- Optics, acoustics, quantum physics, telecommunications, radar, control systems, signal processing, speech recognition, data compression, image processing, seismology, mass spectrometry...
- Digital media. [DVD, JPEG, MP3, H.264]
- Medical diagnostics. [MRI, CT, PET scans, ultrasound]
- Numerical solutions to Poisson's equation.
- Shor's quantum factoring algorithm.

“The FFT is one of the truly great computational developments of [the 20th] century. It has changed the face of science and engineering so much that it is not an exaggeration to say that life as we know it would be very different without the FFT.”

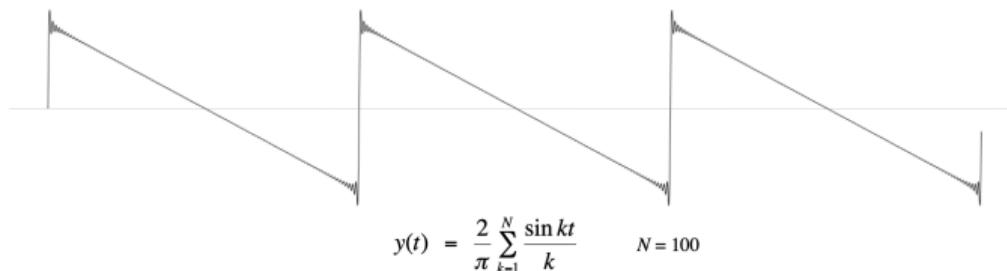
— Charles van Loan



Gilbert Strang described the FFT as “the most important numerical algorithm of our lifetime”.

Fourier Analysis

Fourier theorem. [Fourier, Dirichlet, Riemann] Any (sufficiently smooth) periodic function can be expressed as the sum of a series of sinusoids.



Euler's identity.

$$e^{ix} = \cos x + i \sin x.$$

Sinusoids. Sum of sine and cosines = sum of complex exponentials.

Fourier Transform

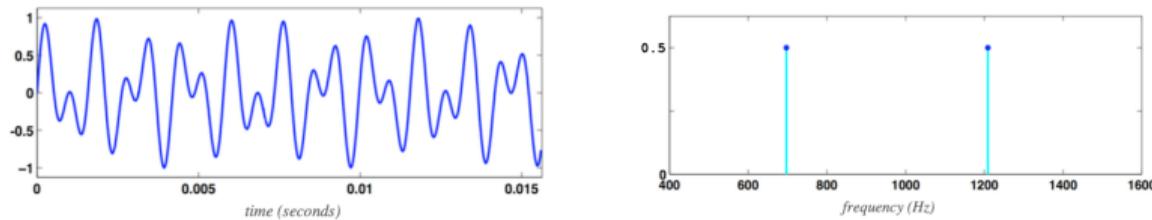


Figure: time domain vs. frequency domain

Signal processing viewpoint of FFT

Fast way to convert between time-domain and frequency-domain

Algorithmic viewpoint of FFT

Fast way to multiply and evaluate polynomials

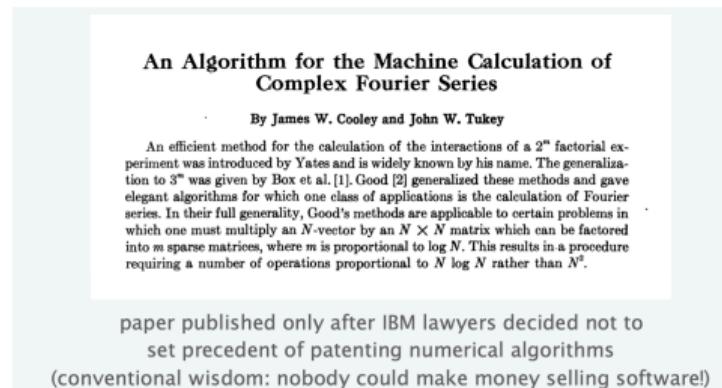
FFT: Brief History

Gauss. Analyzed periodic motion of asteroid Ceres (in Latin)

Runge-König (1924). Laid theoretical groundwork.

Danielson-Lanczos (1942). Efficient algorithm, X-ray crystallography.

Cooley-Tukey (1965). Monitoring nuclear tests in Soviet Union and tracking submarines. Rediscovered and popularized FFT.



Importance not fully realized until the advent of digital computers.

Story of Cooley-Tukey Algorithm

James Cooley and John Tukey independently rediscovered these earlier algorithms^[7] and published a more general FFT in 1965 that is applicable when n is composite and not necessarily a power of 2, as well as analyzing the $O(n \log n)$ scaling.^[13] Tukey came up with the idea during a meeting of President Kennedy's Science Advisory Committee where a discussion topic involved detecting nuclear tests by the Soviet Union by setting up sensors to surround the country from outside. To analyze the output of these sensors, an FFT algorithm would be needed. In discussion with Tukey, Richard Garwin recognized the general applicability of the algorithm not just to national security problems, but also to a wide range of problems including one of immediate interest to him, determining the periodicities of the spin orientations in a 3-D crystal of Helium-3.^[14] Garwin gave Tukey's idea to Cooley (both worked at IBM's Watson labs) for implementation.^[15] Cooley and Tukey published the paper in a relatively short time of six months.^[16] As Tukey did not work at IBM, the patentability of the idea was doubted and the algorithm went into the public domain, which, through the computing revolution of the next decade, made FFT one of the indispensable algorithms in digital signal processing.

FFT in Practice

Fastest Fourier transform in the West. [Frigo and Johnson]

- Optimized C library.
- Features: DFT, DCT, real, complex, any size, any dimension.
- Won Wilkinson Prize '99.

Implementation details.

- Instead of executing predetermined algorithm, it evaluates your hardware and uses a special-purpose compiler to generate an optimized algorithm catered to “shape” of the problem.
- Core algorithm is nonrecursive version of Cooley-Tukey.
- $\Theta(n \log n)$, even for prime sizes.



Integer Multiplication, Redux

Integer multiplication. Given two n bit integers $a = a_{n-1} \dots a_1 a_0$ and $b = b_{n-1} \dots b_1 b_0$, compute their product ab .

- ① Form two polynomials (base-2 representation) $\Rightarrow a = A(2)$, $b = B(2)$

$$A(x) = a_0 + a_1x + \dots + a_{n-1}x^{n-1}$$

$$B(x) = b_0 + b_1x + \dots + b_{n-1}x^{n-1}$$

- ② Compute $C(x) = A(x)B(x)$ via FFT, evaluate $C(2) = ab$

Running time: $\Theta(n \log n)$ complex arithmetic operations.

Practice. GMP uses brute force, Karatsuba, and FFT, depending on the size of n .



Figure: the fastest bignum library on the planet

Summary (1/3)

Concept of Divide-and-Conquer

Main Idea. Reduce problems to subproblems

Principle. Subproblems should be of the same type of the original problem, and can be solved individually.

- Direct dividing: splitting original problem into subproblems with roughly same size
 - FindMinMax, Merge Sort
- Sophisticated dividing
 - General selection: using median of median as pivot (finding the pivot itself requires effort)
 - Closest pair of points: analysis of the strip around the midline
 - Convex hull: sometimes it is hard to split in a balanced manner (convex hull)

Summary (2/3)

Implementation. Recursion or iteration (be careful of the smallest subproblem which can be solved outright)

Time complexity

- Finding the recurrence relation and initial values, solving the recurrence relation

Recurrence relation of divide-and-conquer algorithm

$$T(n) = aT(n/b) + f(n)$$

- a : #(subproblems), n/b : size of subproblems
- $f(n)$: cost of dividing and combining

Summary (3/3)

Optimization trick 1. Reduce the number of subproblems: when $f(n)$ is not very large, $h(n) = n^{\log_b^a}$ dominates the overall complexity $\Rightarrow T(n) = \Theta(h(n))$

- Reduce a can immediately lower the order of $T(n)$
- When subproblems are related \leadsto exploit relations to solve some subproblems by combining the solutions to other subproblems

Examples

- power algorithm: subproblems are same
 - simple algebraic trick: integer multiplication ($f(n)$ is still low)
 - exploit dependence: matrix multiplication ($f(n)$ may increase but does not affect the order)
-

Optimization trick 2. Reduce the cost of dividing and combining $f(n)$: add global preprocessing

- closest pair of points

Important Divide-and-Conquer Algorithms

Searching algorithm: binary search

Sorting algorithm: quick sort, merge sort

Selection algorithm: find min/max, general selection algorithm

Closest pair of points, Convex hull

Fast Power algorithm

Multiplying matrices: Strassen algorithm

Multiplying integers, polynomials: FFT

从前慢

随着算法不断拨快生活的节奏，内心越希望可以偶尔慢下来，松下烹茶、雨夜听琴，
享受一份从容安宁。 ——我的心声

从前慢

随着算法不断拔快生活的节奏，内心越希望可以偶尔慢下来，松下烹茶、雨夜听琴，
享受一份从容安宁。
——我的心声



Figure: 木心



Figure: 云雀叫了一整天

从前慢

随着算法不断拔快生活的节奏，内心越希望可以偶尔慢下来，松下烹茶、雨夜听琴，享受一份从容安宁。
——我的心声

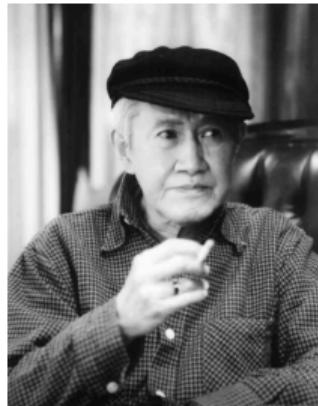


Figure: 木心



Figure: 云雀叫了一整天

从前慢，慢的不仅是车、马与邮件，还有在等待中默默酝酿的心。一封薄薄的信，装着一份炽热的情，翻过山淌过水送至一生唯一的爱人手中，信上的折痕都是那么饱含爱意…

——网易云音乐评论